Top 30 AI, ML & GenAI Interview Questions

Open in app ↗

Medium

7 min read · Dec 20, 2025

Robi Kumar Tomar Follow

Listen Share ••• More

After each interview, I jot down the questions and discussion points in a notebook while they're still fresh — that habit is how this list of 30 questions came together.

AI/ML interviews have evolved — yet many candidates are still approaching them the old way. They're no longer about recalling definitions or reciting algorithms. Today, interviewers are quietly trying to answer one critical question:

> **"Can this person actually build, evaluate, and own an AI/ML system in the real world?"**

After participating in and observing multiple AI/ML interviews across companies and roles, one pattern has become impossible to ignore:
Most candidates don't fail because they lack knowledge; they fail because they lack clarity, intuition, and production-level awareness.

This article distills the 30 AI, ML, and GenAI interview questions that repeatedly came up in my recent interviews — not as rote Q&A, but as thinking frameworks.
If you truly understand these, you won't just blend in with the crowd — you'll stand out as someone interviewers trust to deliver real-world AI/ML systems.

1. How Do You Detect and Fix Overfitting?

Symptoms

- Low training error, high validation error

Fixes

- Regularization
- Simpler model
- More data
- Feature reduction
- Better cross-validation

What interviewers look for:
Not just techniques — but your ability to **diagnose the cause**, not blindly apply fixes.

2. L1 vs L2 Regularization — When Would You Use Each?

- **L1 (Lasso)**: induces sparsity → implicit feature selection
- **L2 (Ridge)**: smooths weights → stability under noisy data

Practical insight

> If model interpretability or feature pruning matters, L1 is useful. If features are correlated and stability matters, L2 is safer.

3. What Is the Bias–Variance Trade-Off?

Bias–variance describes the balance between:

- **Bias**: error from overly simplistic assumptions (underfitting)
- **Variance**: error from excessive sensitivity to data noise (overfitting)

The goal is not minimizing either independently, but minimizing **generalization error**.

Interview-ready insight

> In fast-changing domains (fraud, pricing, user behavior), some variance is acceptable if retraining is frequent. In stable domains, lower variance models are preferred.

4. What Is Cost-Aware Machine Learning?

The best model is not always the most accurate.

Real constraints include:

- Inference latency
- Compute cost
- Infrastructure limits

Interviewers value candidates who balance **accuracy with feasibility**.

5. What Is a Confusion Matrix and Why Is It Important?

A confusion matrix reveals **error types**, not just error counts.

Interviewers expect you to explain how it:

- Guides metric selection
- Aligns model optimization with business cost

6. Parametric vs Non-Parametric Models

- **Parametric**: fixed number of parameters (Linear Regression)
- **Non-parametric**: complexity grows with data (kNN, trees)

Key trade-off

Interpretability vs flexibility — and how dataset size influences the decision.

7. What Is Cross-Validation and When Can It Fail?

Cross-validation estimates generalization performance.

Critical pitfall (often missed):

- Random splits **must not** be used for time-series or sequential data.

Strong answer

> "I use time-aware splits to avoid future information leakage."

8. What Is Data Leakage?

Data leakage occurs when information unavailable at prediction time influences training.

Examples:

- Using future data in feature engineering
- Leaking target-derived features

Why interviewers care:
Leakage produces *great offline metrics and terrible production models*.

9. Feature Engineering — Why Is It More Important Than Models?

Most real performance gains come from **better features**, not better algorithms.

Example
In delivery ETA prediction:

- Driver reliability trends
- Warehouse congestion signals
- Recent delay volatility
  often outperform model complexity improvements.

10. How Do You Handle Missing Data?

Options depend on:

- Missingness pattern
- Feature importance

- Business impact

Techniques include:

- Dropping rows
- Statistical imputation
- Model-based imputation

Good answers explain **why**, not just **how**.

11. Bagging vs Boosting

- **Bagging** reduces variance (Random Forest)
- **Boosting** reduces bias (XGBoost, AdaBoost)

Interview insight

> Boosting can overfit noisy data — monitoring validation trends is critical.

12. What Is ROC–AUC and When Is It Misleading?

ROC–AUC measures class separability across thresholds.

Limitation

- Can appear strong even when precision/recall at operating thresholds is poor — especially with imbalanced data.

13. Feature Importance — How Do You Explain Models?

Methods include:

- Tree-based importance
- Permutation importance
- SHAP values

What interviewers want:
Your ability to translate explanations into **business decisions**, not plots.

14. Hyperparameter Tuning Strategies

- Grid search

- Random search
- Bayesian optimization

Trade-off: **search cost vs marginal performance gain**.

15. What Happens After Model Training?

This is where many candidates fail.

Real systems require:

- Deployment strategy
- Monitoring
- Retraining triggers
- Rollback mechanisms

Key failure mode

> **Training–serving skew** — features computed differently offline vs online.

16. What Is Data Drift vs Concept Drift?

- **Data drift**: input distribution changes
- **Concept drift**: relationship between inputs and outputs changes

Both silently degrade models if unmonitored.

17. How Do You Monitor Models in Production?

Common signals:

- Feature distribution shifts
- Prediction confidence changes
- Business KPI degradation

Strong answers emphasize **early detection**, not reactive fixes.

18. Precision, Recall, and F1 — How Do You Choose?

- **Precision**: correctness of positive predictions

- **Recall:** coverage of actual positives
- **F1:** balance between both

Non-generic example
In a **logistics delay prediction system**, false alarms may overload operations teams, while missed delays damage SLAs. The "best" metric depends on operational cost — not theory.

19. Deep Learning Essentials You Should Explain Clearly

You should comfortably explain:

- CNNs vs RNNs
- Backpropagation
- Activation functions
- Gradient descent variants

Clarity matters more than equations.

20. Generative AI & LLMs — What's Expected Now?

Even non-GenAI roles expect awareness of:

- Transformer architecture
- Prompting vs fine-tuning
- Hallucination risks

Crucial insight

> LLM evaluation is fundamentally harder — often requiring human review or task-specific success metrics instead of accuracy.

21. What Is RAG (Retrieval-Augmented Generation) and Why Is It Needed?

RAG combines:

- **Retrieval** (fetching relevant documents)
- **Generation** (LLM response synthesis)

It solves a core LLM limitation:

👉 *LLMs do not have up-to-date or domain-specific knowledge.*

Why interviewers ask this

They want to know if you understand **how GenAI systems work beyond prompts.**

Strong insight

> "RAG reduces hallucination risk and avoids frequent fine-tuning by grounding responses in external knowledge."

22. How Do You Prevent Hallucinations in LLMs?

Hallucinations are confident but incorrect outputs.

Mitigation strategies

- Retrieval grounding (RAG)
- Constrained prompting
- Output validation rules
- Human-in-the-loop review
- Confidence thresholds

Interview-grade answer

> "Hallucinations can't be fully eliminated — only managed through system design."

That sentence signals **maturity.**

23. Prompt Engineering vs Fine-Tuning — When Do You Choose Each?

```
| Scenario                  | Better Choice                      |
| ------------------------- | ---------------------------------  |
| Rapid iteration           | Prompt engineering                 |
| Stable domain language    | Fine-tuning                        |
| Sensitive enterprise data | RAG + prompts                      |
| Cost-sensitive inference  | Fine-tuning or smaller models      |
```

What interviewers test
Your ability to balance **cost, flexibility, and maintainability**.

24. What Is Tokenization and Why Does It Matter?

Tokenization splits text into subword units used by LLMs.

Why it matters

- Cost scales with token count
- Context window limits
- Prompt truncation risks

Practical example

> "Two prompts with the same text length can have very different token costs."

This shows **operational awareness**, not theory.

25. How Do You Evaluate LLM Outputs?

This is one of the **most asked GenAI questions today**.

Key challenge
There is no single "accuracy" metric.

Evaluation methods

- Human evaluation
- Task-specific heuristics
- Reference-based scoring
- Downstream business metrics

Strong statement

> "LLM evaluation is probabilistic and contextual — not deterministic."

26. What Are Embeddings and How Are They Used?

Embeddings convert text into dense vectors that capture semantic meaning.

Common use cases

- Semantic search
- Clustering
- Recommendation
- Retrieval for RAG

Why interviewers care
Embeddings are foundational to **most real GenAI systems**, not just chatbots.

27. How Do Context Windows Affect LLM Design?

Context window = maximum tokens a model can process at once.

Implications

- Long documents must be chunked
- Poor chunking causes loss of meaning
- Retrieval ranking becomes critical

Interview insight

> "More context isn't always better — relevance matters more than length."

28. What Is Function Calling / Tool Use in LLMs?

LLMs can:

- Call APIs
- Execute tools
- Trigger workflows

Why it matters
This turns LLMs from *text generators* into **agents**.

Example

> "An LLM extracts intent, calls a pricing API, then summarizes the result."

This signals **system-level thinking.**

29. What Are the Security Risks in GenAI Systems?

High-risk areas:

- Prompt injection
- Data leakage
- Unauthorized tool execution
- Training data memorization

Interview-ready answer

> "LLMs expand the attack surface — security must be designed, not assumed."

This is increasingly asked in enterprise interviews.

30. What Is the Difference Between an LLM and an AI Agent?

- **LLM**: generates text based on probability
- **Agent**: plans, decides, executes, and observes outcomes using tools

Why interviewers ask
They want to know if you understand **modern agentic architectures**, not just chat interfaces.

Common AI/ML Interview Mistakes

- Giving definitions without intuition
- Ignoring data quality and edge cases
- Treating deployment as "out of scope"
- Naming tools instead of explaining principles
- Optimizing metrics without understanding business impact

Avoid these mistakes — they cost offers.

Final Interview Strategy

When answering **any** AI interview question:

1. Clarify the problem
2. Explain the intuition
3. Discuss trade-offs
4. Tie the decision to real-world impact
5. Summarize clearly

This structure consistently signals **senior-level thinking**.

Wrapping Up

Modern AI/ML interviews don't test how much you remember — they test **how well you reason under uncertainty**.

If you truly understand these **30 topics**, you won't just answer questions — you'll earn **trust**.

And in AI hiring, **trust is everything**.

❤️ **Robi Tomar here — thanks for reading till the end.**

If this sparked a new way of thinking, hit that clap 👏 and follow for more practical AI engineering insights.

Artificial Intelligence | Machine Learning | Data Science | Software Development | Software Engineering

Follow

Written by Robi Kumar Tomar

357 followers · 26 following

Architect | 15+ yrs exp | Scalable systems & applied AI. A love for learning, building & sharing.

No responses yet

Kun Ma

What are your thoughts?

More from Robi Kumar Tomar

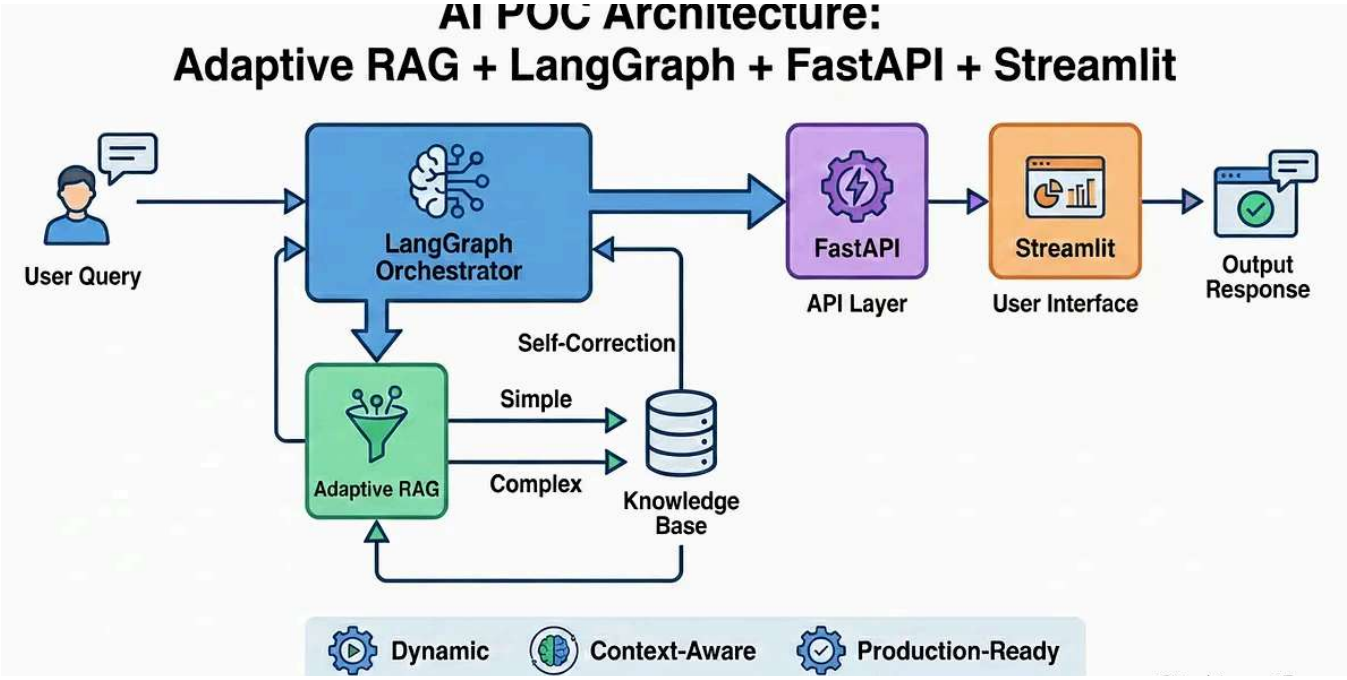

Robi Kumar Tomar

Build a Full POC Using Adaptive RAG + LangGraph + FastAPI + Streamlit (Complete Step-by-Step Guide)

AI applications today must go beyond simple prompt-response systems. They need to be dynamic, context-aware, and production-ready.

Dec 13, 2025 112 2

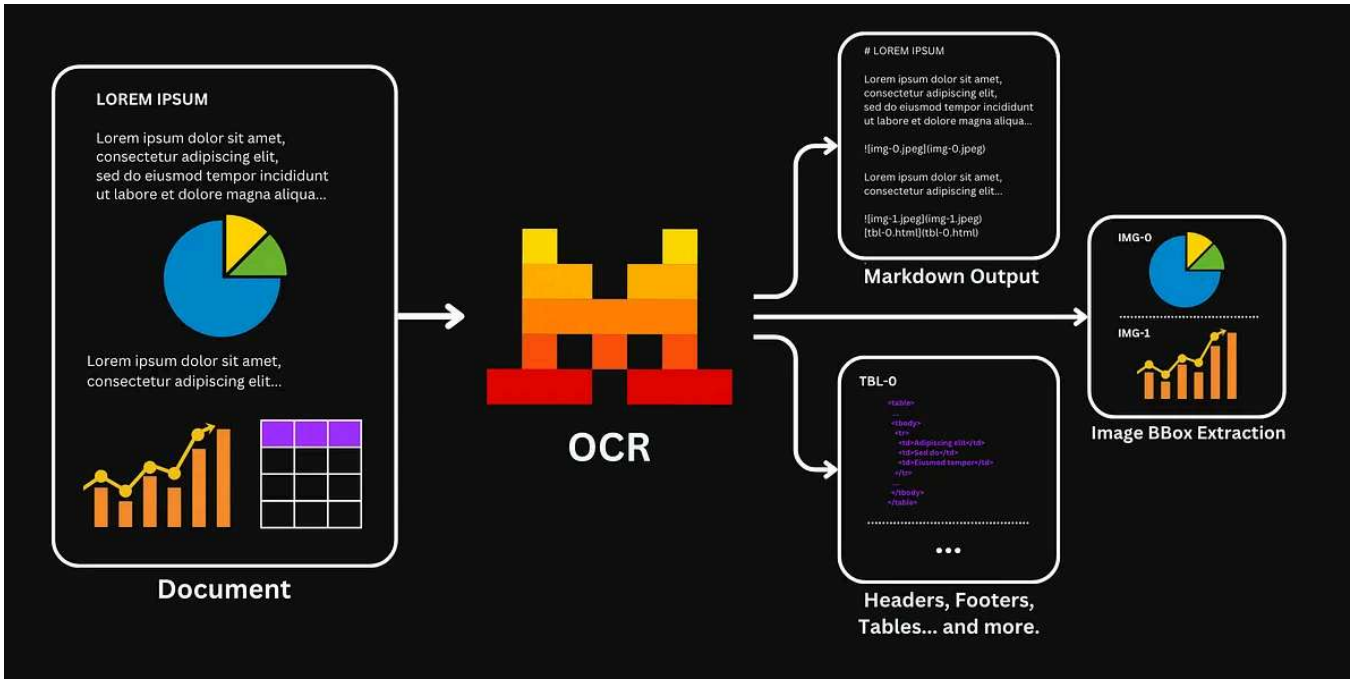

In CodeToDeploy by Robi Kumar Tomar

Mistral OCR 3 Changed OCR Overnight—It's Now AI Infrastructure

For most of my career, OCR was the boring part of the stack—necessary plumbing you bolted on so smarter systems could do their work.

Dec 25, 2025 275 2

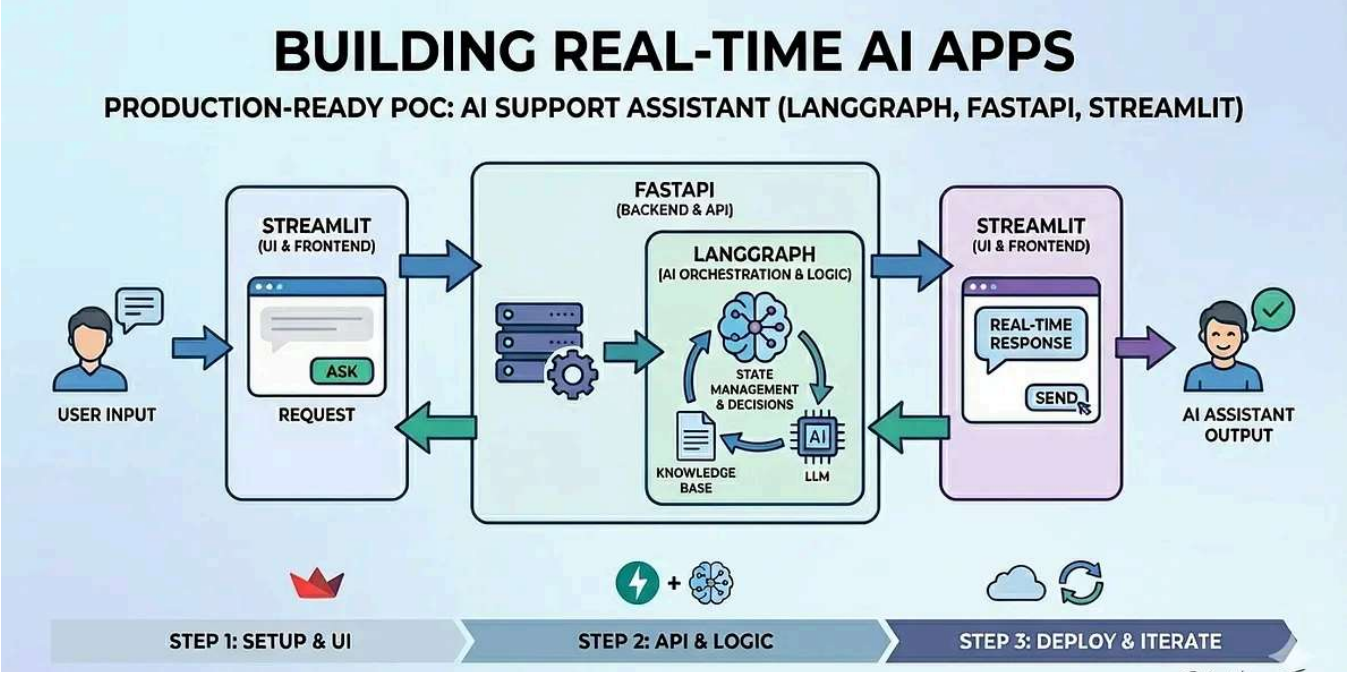

Robi Kumar Tomar

Building Real-Time AI Apps with LangGraph, FastAPI & Streamlit

A Production-Ready Proof-of-Concept (Step-by-Step)

Dec 19, 2025 44 1

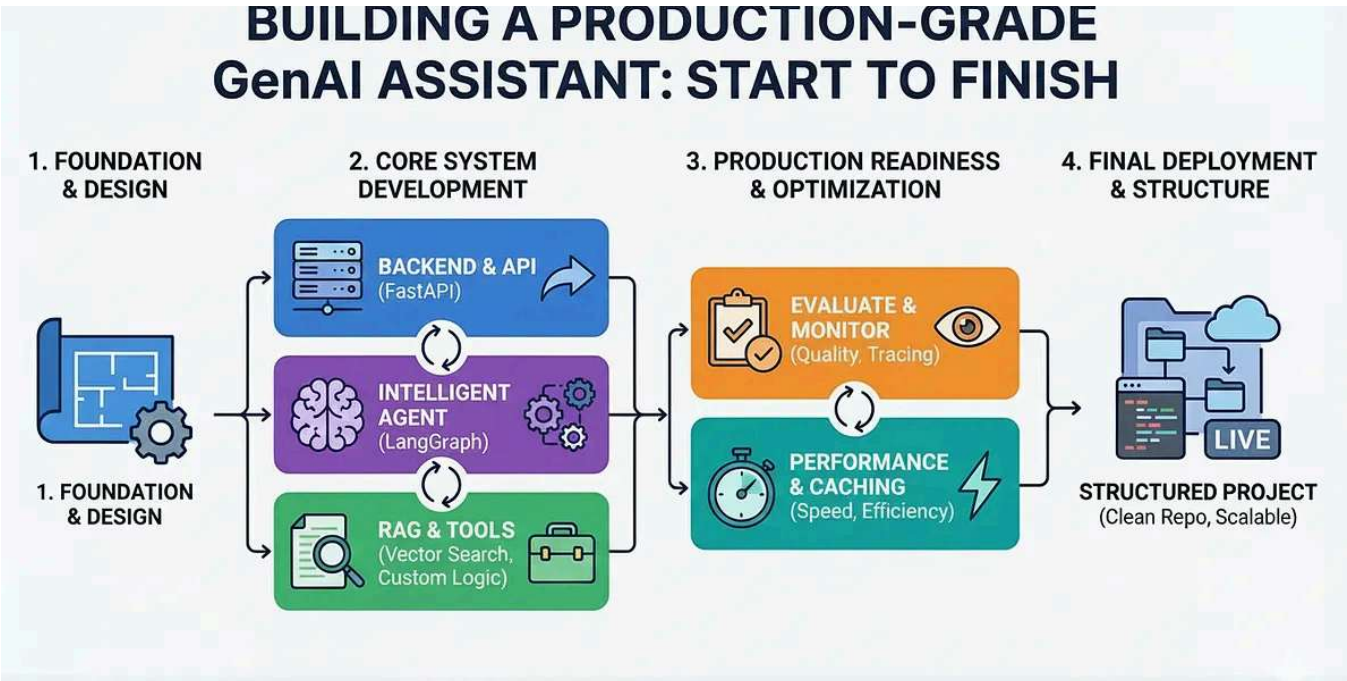

In AI Mind by Robi Kumar Tomar

Build a Real End-to-End GenAI Project : Agents + RAG + Tools + Evaluations + Monitoring (Complete…

Generative AI Series—Part 24

Dec 2, 2025 202 3

See all from Robi Kumar Tomar

Recommended from Medium

In Level Up Coding by Christian Bernecker

Building an AI Agent from Scratch with pure Python

From Theory to Implementation: Building a Robust, Self-Correcting AI Agent from Scratch with Python

Feb 17 474 4

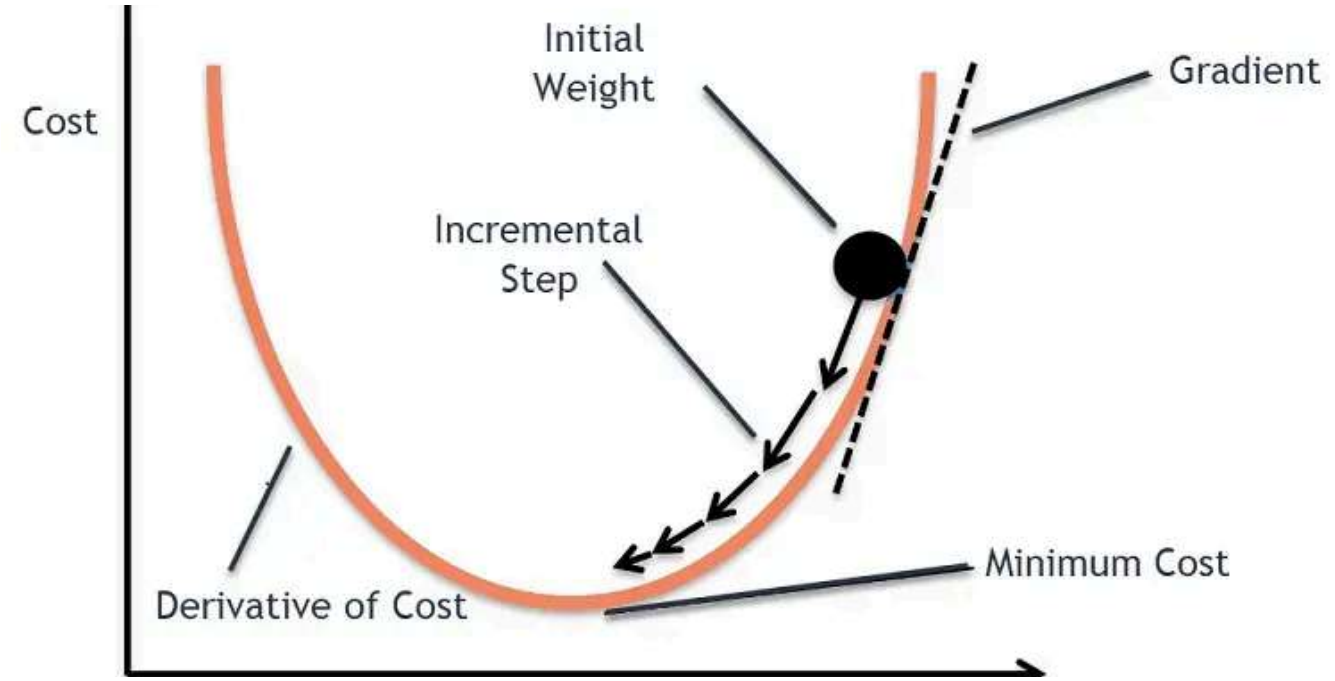

MOHAMMED MOSEENA

My ML Interview Notebook: A Revision Guide for Data Science Interviews

ML Revision Notes

Nov 19, 2025 413 12

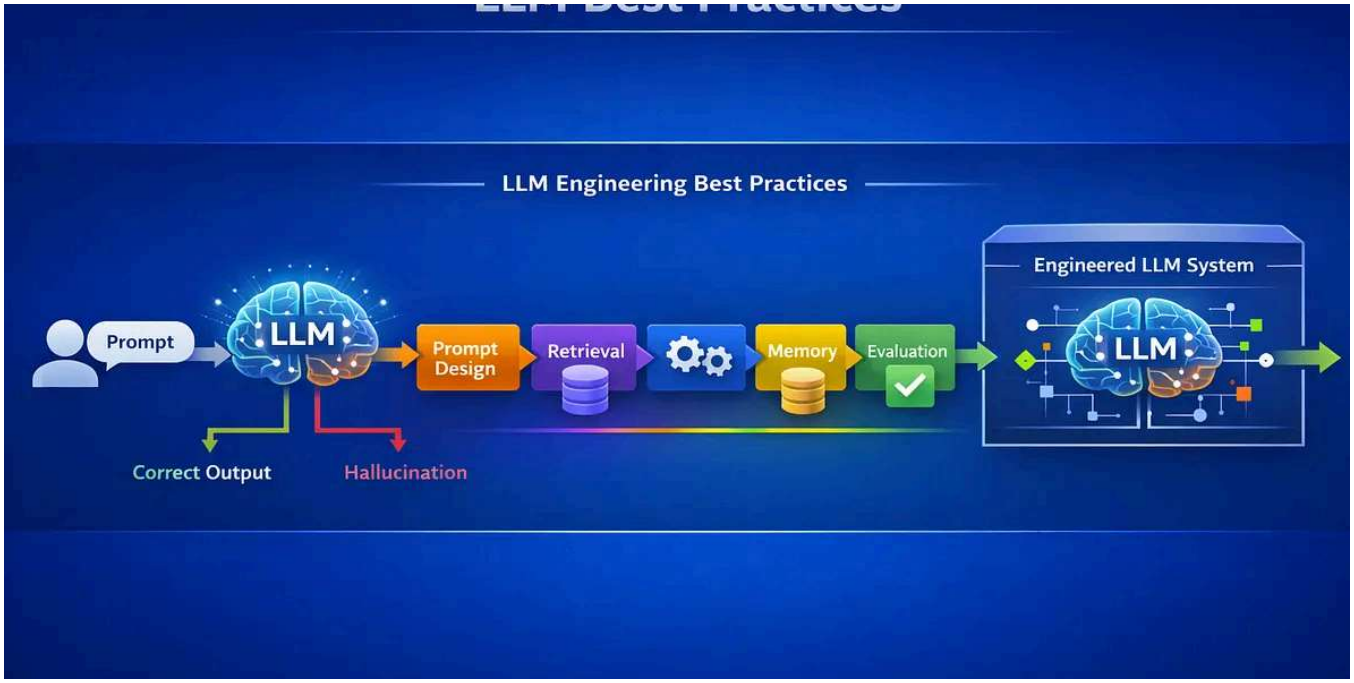

In Towards AI by Khushbu Shah

25 LLM Best Practices I Believe Every Engineer Should Learn Early

Anyone can get an LLM to respond. Making it respond correctly every time is engineering. These LLM best practices will help you turn models...

Feb 16 286 2

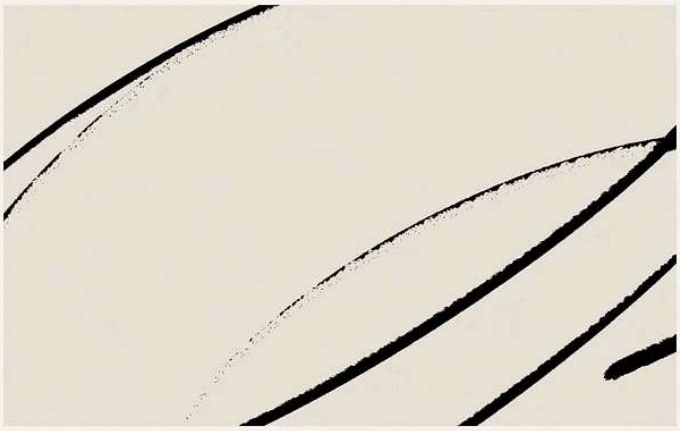

Simplilearn

60+ Machine Learning Interview Questions and Answers

TL;DR: Need machine learning interview questions you can actually use in a hiring loop? This article covers 60+ machine learning interview...

Jan 14 22

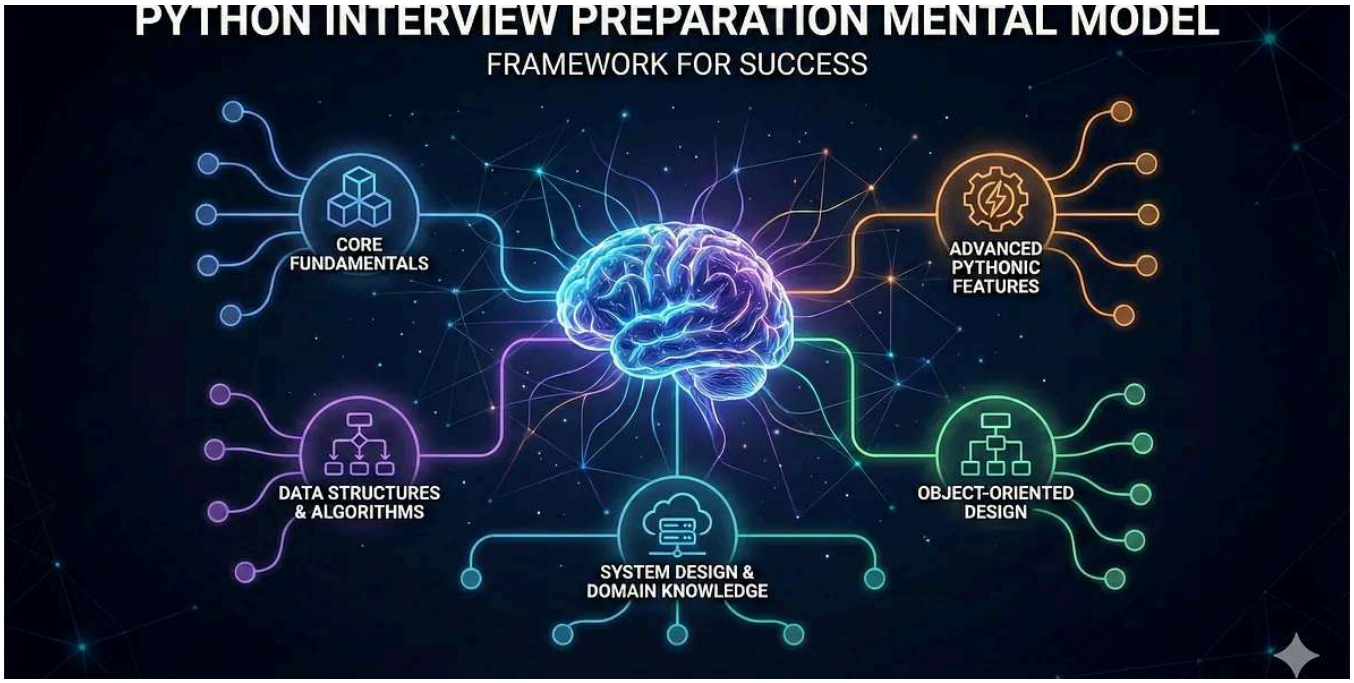

In Activated Thinker by Shreyas Naphad

Python Interview Preparation

A complete beginner-to-intermediate Python revision guide for data, ML, and analyst interviews.

Jan 23 86

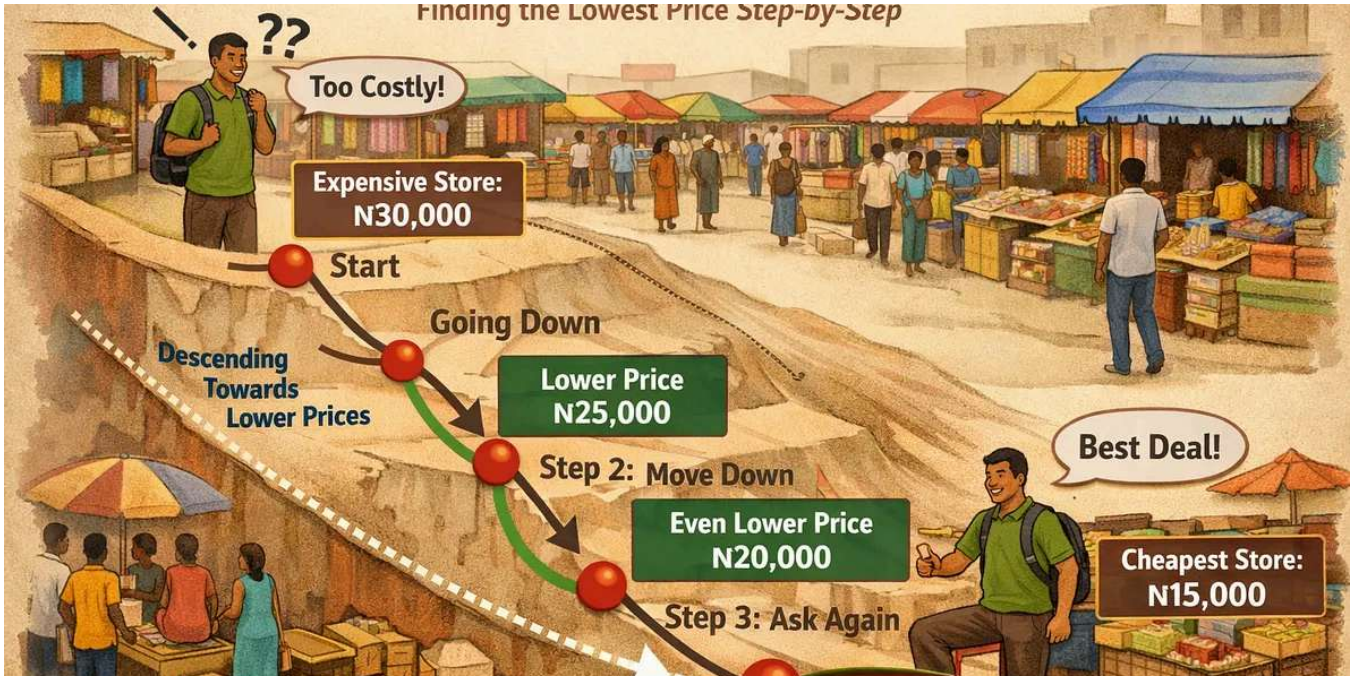

Jeremiah J. U. Omueti

Understanding Logistic Regression: A Beginner's Guide to Classification

Logistic Regression is really about figuring out how to classify a thing. Am I looking at a horse, a zebra, or a unicorn; or in the case of...

Jan 14 24

See more recommendations